

## Project 2 : In the Interest of Time

*Due ~~10/21/2015~~ 2/28/2023 @ 12pm (noon)*

### TOPICS:

- ↳ NVIDIA Threads Architecture
- ↳ Time Library
- ↳ File I/O
- ↳ Random Numbers
- ↳ Parallel Algorithm Analysis



### BACKGROUND:

During class, we discussed how to implement a matrix multiplication subroutine. We also discussed several metrics for measuring parallel algorithm performance.

### PART 1: MATRIX MULTIPLY

Design and implement C and CUDA programs that called MatrixMult.c, MatrixMultNaive.cu, and MatrixMultTiled.cu.

The MatrixMult.c, MatrixMultNaive.cu, and MatrixMultTiled.cu programs will use a command line argument that corresponds to the row/column length to multiply two square matrices with random single precision numbers (i.e., floats). They will then output 1) the entire execution time in seconds from the beginning to the end of the function/kernel call only, and 2) the product of the two matrices with each element printed out to a file called "product.dat" in a tab-delimited, row/column format. The C program must implement a serial code via a function call. The CUDA MatrixMultNaive.cu program must implement a parallel code via a kernel call using just global memory, and the MatrixMultTiled.cu program must implement a parallel code via a kernel call using tiles and shared memory.

You must check for 1) the appropriate number of command line arguments and 2) whether the command line argument corresponds to a positive number (you may round floating point inputs). Appropriate error messages must be issued, followed by a graceful exit.

To analyze the relative performances of the CPU vs. GPU implementations of your matrix multiplication algorithms, you will compute the time it takes to complete a matrix multiplication function/kernel call. For the GPU implementations, you will show how the block sizes impact performance by running them with tile sizes of

16x16 and 32x32 (you can leave your final code with either tile size). You will then construct graphs (CPU vs. Naive and CPU vs. Tiled) of the speedup vs. matrix size (i.e., row/column length) for several tile sizes to see how their performances vary with system size and tile size. It is up to you to choose appropriate values for the matrix size to construct your graphs, but a matrix sizes of 1,021 and 2,015 will be used for testing during grading.

Hint: In the execution of your GPU implementation, for accurate timing information, you must call the kernel twice because the GPU initialization takes time that is unrelated to the algorithm.

## PART 2: IMAGE BLUR

You are given a working program that blurred an image. Apply the program to several bmp images of different sizes and graph how its performance varies with system size. Modify the program using shared memory to improve its performance and compare how its performance varies with system size as compared to the original program.

Write a *brief* description of the strategy you employed to improve the performance of the blur program. You will not be graded on how much your strategy improved the performance or even whether performance improved. All I am looking for here is that you gave it some thought and you were able to reason out a strategy.

## IMPLEMENTATION NOTES:

Any program that does not compile or does not have correctly constructed Makefiles will not be graded. The C program must compile using g++ and the CUDA program must compile using nvcc.

## COMMENTS AND STYLE:

Although there will be no formal policy on commenting and style, the reader should be able to easily follow the main purpose of the code. Each set of code that does something significant must be commented. The variable names should be easily recognizable and acronyms should be avoided if possible.

*Do not be surprised if help is not forthcoming if your code is poorly commented and/or difficult to follow. You have been warned.*

## PROJECT SUBMISSION:

Your project will be submitted via Github. The invitation to your repository is here:

Part 1: <https://classroom.github.com/a/5h9vo67e>

Part 2: [https://classroom.github.com/a/RHZgJ\\_SQ](https://classroom.github.com/a/RHZgJ_SQ)

## PLEDGED WORK POLICY:

Assignments in Computer Science courses may be specified as "pledged work" assignments by the professor of the course. When an assignment is specified as "pledged work" the only aid that the student may seek is from either the course

professor or an assistant that the professor has explicitly specified. On "pledged work" assignments the student may not use the services of a tutor.

For this project, you may discuss only basic C/C++ syntax with others. Any other discussions of the project are strictly prohibited except with the professor of the course. Your code and your implementation of the project must be the product of your own work.